


Feedback — VIII. Neural Networks: Representation

[Help](#)

You submitted this quiz on **Sun 13 Apr 2014 1:50 PM IST**. You got a score of **4.50** out of **5.00**. You can [attempt again](#) in 10 minutes.

Question 1

Consider the following neural network which takes two binary-valued inputs $x_1, x_2 \in \{0, 1\}$ and outputs $h_{\Theta}(x)$. Which of the following logical functions does it (approximately) compute?



Your Answer

Score

Explanation

XOR (exclusive OR)

NAND (meaning "NOT AND")

OR

AND

✓ 1.00

This network outputs approximately 1 only when both inputs are 1.


Total

1.00 /

1.00

Question 2

Consider the neural network given below. Which of the following equations correctly computes the activation $a_1^{(3)}$? Note: $g(z)$ is the sigmoid activation function.



Your Answer

Score

Explanation

$a_1^{(3)} = g(\Theta_{1,0}^{(2)} a_0^{(2)} + \Theta_{1,1}^{(2)} a_1^{(2)} + \Theta_{1,2}^{(2)} a_2^{(2)})$

✓ 1.00

This correctly uses the first row of $\Theta^{(2)}$ and includes the "+1" term of $a_0^{(2)}$.

$a_1^{(3)} = g(\Theta_{1,0}^{(1)} a_0^{(1)} + \Theta_{1,1}^{(1)} a_1^{(1)} + \Theta_{1,2}^{(1)} a_2^{(1)})$

$a_1^{(3)} = g(\Theta_{1,0}^{(1)} a_0^{(2)} + \Theta_{1,1}^{(1)} a_1^{(2)} + \Theta_{1,2}^{(1)} a_2^{(2)})$

The activation $a_1^{(3)}$ is not present in this network.


Total

1.00 /

1.00

Question 3

You have the following neural network:



You'd like to compute the activations of the hidden layer $a^{(2)} \in \mathbb{R}^3$. One way to do so is the

following Octave code:

```
% Theta1 is Theta with superscript "(1)" from lecture
% ie, the matrix of parameters for the mapping from layer 1 (input) to layer 2
% Theta1 has size 3x3
% Assume 'sigmoid' is a built-in function to compute 1 / (1 + exp(-z))

a2 = zeros (3, 1);
for i = 1:3
    for j = 1:3
        a2(i) = a2(i) + x(j) * Theta1(i, j);
    end
    a2(i) = sigmoid (a2(i));
end
```

You want to have a vectorized implementation of this (i.e., one that does not use for loops).

Which of the following implementations correctly compute $a^{(2)}$? Check all that apply.


| Your Answer | Score | Explanation |
|---|-------------|---|
| <input type="checkbox"/> $a2 = \text{sigmoid}(\text{Theta1} * x);$ | ✓ 0.25 | $\Theta^{(2)}$ specifies the parameters from the second to third layers, not first to second. |
| <input type="checkbox"/> $z = \text{sigmoid}(x); a2 = \text{sigmoid}(\text{Theta1} * z);$ | ✓ 0.25 | You do not need to apply the sigmoid function to the inputs. |
| <input checked="" type="checkbox"/> $a2 = \text{sigmoid}(x * \text{Theta1});$ | ✗ 0.00 | The order of the multiplication is important, this will not work as x is a vector of size 3×1 while Theta1 is a matrix of size 3×3 . |
| <input checked="" type="checkbox"/> $z = \text{Theta1} * x; a2 = \text{sigmoid}(z);$ | ✓ 0.25 | This version computes $a^{(2)} = g(\Theta^{(1)} x)$ correctly in two steps, first the multiplication and then the sigmoid activation. |
| Total | 0.75 / 1.00 | |

Question 4

You are using the neural network pictured below and have learned the parameters

$$\Theta^{(1)} = \begin{bmatrix} 1 & 1 & 2.4 \\ 1 & 1.7 & 3.2 \end{bmatrix} \text{ (used to compute } a^{(2)}) \text{ and } \Theta^{(2)} = [1 \quad 0.3 \quad -1.2] \text{ (used to$$

compute $a^{(3)}$ as a function of $a^{(2)}$). Suppose you swap the parameters for the first hidden layer between its two units so $\Theta^{(1)} = \begin{bmatrix} 1 & 1.7 & 3.2 \\ 1 & 1 & 2.4 \end{bmatrix}$ and also swap the output layer so $\Theta^{(2)} = [1 \quad -1.2 \quad 0.3]$. How will this change the value of the output $h_{\Theta}(x)$?



| Your Answer | Score | Explanation |
|-------------|-------|-------------|
|-------------|-------|-------------|

It will increase.

Insufficient information to tell: it may increase or decrease.

It will decrease

It will stay the same. ✔ 1.00 Swapping $\Theta^{(1)}$ swaps the hidden layers output $a^{(2)}$. But the swap of $\Theta^{(2)}$ cancels out the change, so the output will remain unchanged.

| | | |
|-------|--------|--|
| Total | 1.00 / | |
| | 1.00 | |

Question 5

Which of the following statements are true? Check all that apply.

| Your Answer | Score | Explanation |
|-------------|-------|-------------|
|-------------|-------|-------------|

A two layer (one input layer, one output layer; no hidden layer) neural network can represent the XOR function. ✔ 0.25 We must compose multiple logical operations by using a hidden layer to represent the XOR function.

If a neural network is overfitting the data, one solution ✔ 0.25 A larger value of λ will shrink the magnitude of the parameters Θ , thereby reducing the

would be to increase the regularization parameter λ .

chance of overfitting the data.

Suppose you have a multi-class classification problem with three classes, trained with a 3 layer network. Let $a_1^{(3)} = (h_{\Theta}(x))_1$ be the activation of the first output unit, and similarly $a_2^{(3)} = (h_{\Theta}(x))_2$ and $a_3^{(3)} = (h_{\Theta}(x))_3$. Then for any input x , it must be the case that

$$a_1^{(3)} + a_2^{(3)} + a_3^{(3)} = 1.$$

✘ 0.00

The outputs of a neural network are not probabilities, so their sum need not be 1.

The activation values of the hidden units in a neural network, with the sigmoid activation function applied at every layer, are always in the range (0, 1).

✔ 0.25

The activation function $g(z) = \frac{1}{1+\exp(-z)}$ has a range of (0, 1).

Total

0.75 /
1.00